MONITORING AND LOGGING MANAGEMENT SYSTEM – ONE STOP SOLUTION

By

Vaishal Shah

INTRODUCTION

This is just a demo project I made to showcase my DevOps skills. I integrated different tools and technologies to make things more efficient and less time-consuming.

ELK stack is deployed to make a one-stop solution for monitoring and logging management. ELK stack is deployed using various DevOps technologies. Tech Stack contains; Terraform, Docker, Ansible, GitHub, AWS cloud, and CI/CD. I have used these tools to take the benefits of each technology and make it easier for tech savvies to deploy monitoring and logging infrastructure.

TECHNOLOGY STACK

GitHub - ELK : Setup done.
Terraform - Done. Infrastructure as a code.
Docker - Done. Containerized application.
Ansible - In progress. Configuration management for ELK.
AWS Cloud - Done. Resources are deployed on the AWS cloud.
ELK stack - Done. Tested thoroughly.
CI/CD - In near future (within a week).

GITHUB:

GitHub is the main software I have used as a versioning control system for this project. It consists of an entire repository of the ELK architecture.

It is used for enabling the code version and for continuous integration purposes. People can start contributing to my repository in the future for enhancement.

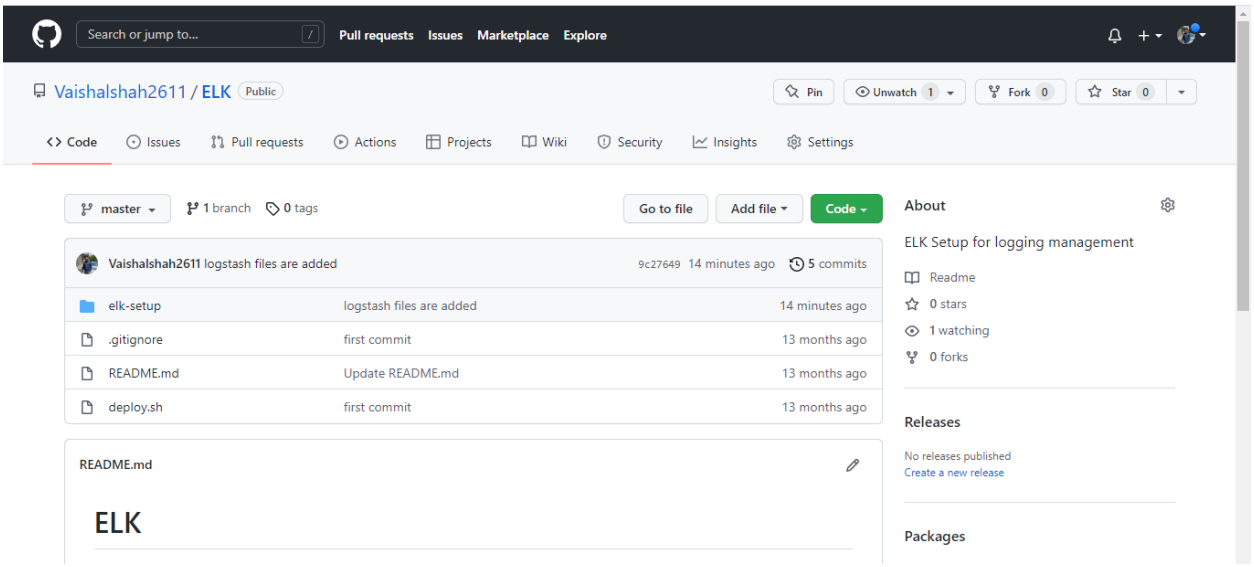

TERRAFORM:

Terraform is used to enable Infrastructure as Code. The infrastructure deployed on AWS using Terraform. Currently, terraform modules are not being used as it becomes easy to pull the defined resources. So, as a beginner, I insisted on writing the code from scratch so that it will help me to clear concepts.

The below screenshot shows the code directory, creation, and deletion of the resources using terraform:

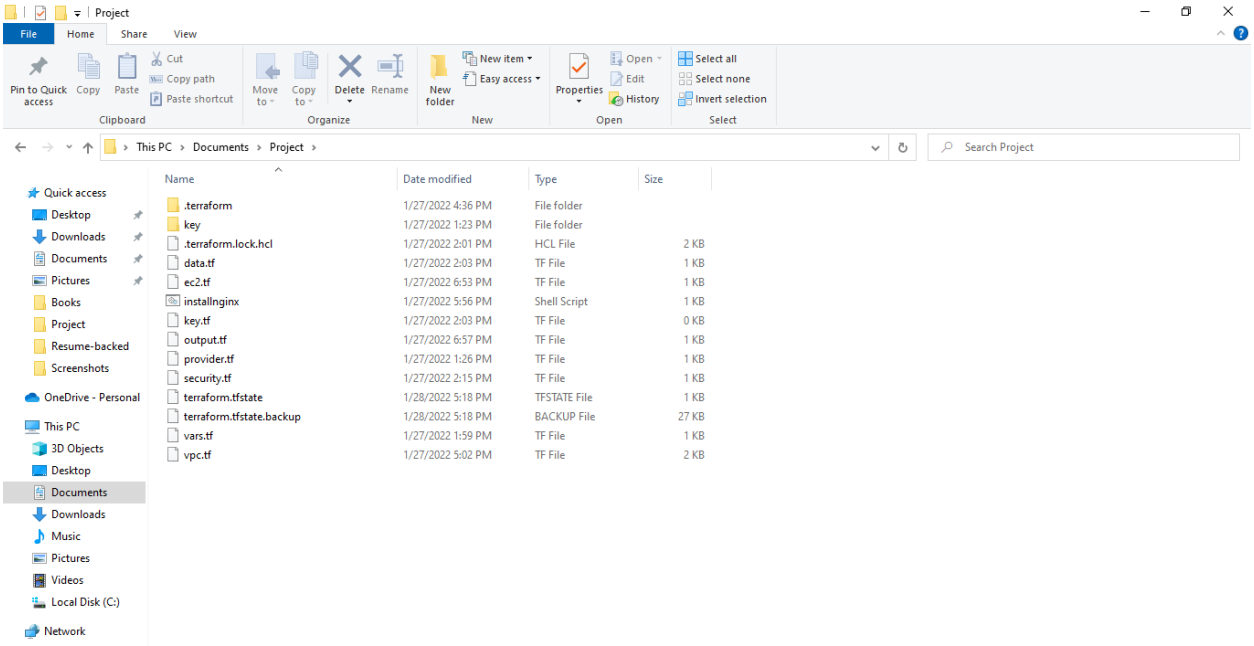
Project
File
Home
Share
View
Pin to Quick access
Copy
Paste
Cut
Copy path
Paste shortcut
Clipboard
Move to
Copy to
Delete
Rename
Organize
New folder
New item
Easy access
New
Properties
Open
Edit
History
Open
Select all
Select none
Invert selection
Select
This PC › Documents › Project ›
Search Project
Quick access
Desktop
Downloads
Documents
Pictures
Books
Project
Resume-backed
Screenshots
OneDrive - Personal
This PC
3D Objects
Desktop
Documents
Downloads
Music
Pictures
Videos
Local Disk (C:)
Network
Name
Date modified
Type
Size
.terraform 1/27/2022 4:36 PM File folder
key 1/27/2022 1:23 PM File folder
.terraform.lock.hcl 1/27/2022 2:01 PM HCL File 2 KB
data.tf 1/27/2022 2:03 PM TF File 1 KB
ec2.tf 1/27/2022 6:53 PM TF File 1 KB
installnginx 1/27/2022 5:56 PM Shell Script 1 KB
key.tf 1/27/2022 2:03 PM TF File 0 KB
output.tf 1/27/2022 6:57 PM TF File 1 KB
provider.tf 1/27/2022 1:26 PM TF File 1 KB
security.tf 1/27/2022 2:15 PM TF File 1 KB
terraform.tfstate 1/28/2022 5:18 PM TFSTATE File 1 KB
terraform.tfstate.backup 1/28/2022 5:18 PM BACKUP File 27 KB
vars.tf 1/27/2022 1:59 PM TF File 1 KB
vpc.tf 1/27/2022 5:02 PM TF File 2 KB

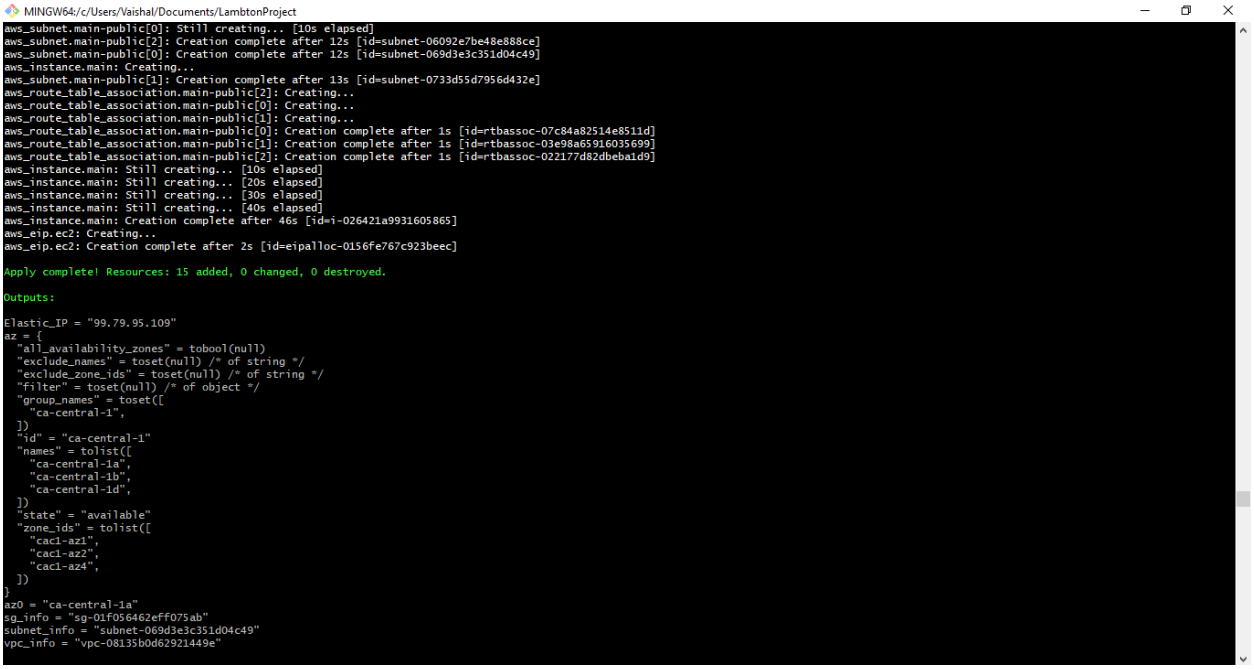
MINGW64:/c/Users/Vaishal/Documents/LambtonProject
aws_subnet.main-public[0]: Still creating... [10s elapsed]
aws_subnet.main-public[2]: Creation complete after 12s [id=subnet-06092e7be48e888ce]
aws_subnet.main-public[0]: Creation complete after 12s [id=subnet-069d3e3c351d04c49]
aws_instance.main: Creating...
aws_subnet.main-public[1]: Creation complete after 13s [id=subnet-0733d55d7956d432e]
aws_route_table_association.main-public[2]: Creating...
aws_route_table_association.main-public[0]: Creating...
aws_route_table_association.main-public[1]: Creating...
aws_route_table_association.main-public[0]: Creation complete after 1s [id=rtbassoc-07c84a82514e8511d]
aws_route_table_association.main-public[1]: Creation complete after 1s [id=rtbassoc-03e98a65916035699]
aws_route_table_association.main-public[2]: Creation complete after 1s [id=rtbassoc-022177d82dbeba1d9]
aws_instance.main: Still creating... [10s elapsed]
aws_instance.main: Still creating... [20s elapsed]
aws_instance.main: Still creating... [30s elapsed]
aws_instance.main: Still creating... [40s elapsed]
aws_instance.main: Creation complete after 46s [id=i-026421a9931605865]
aws_eip.ec2: Creating...
aws_eip.ec2: Creation complete after 2s [id=eipalloc-0156fe767c923beec]
Apply complete! Resources: 15 added, 0 changed, 0 destroyed.
Outputs:
Elastic_IP = "99.79.95.109"
az = {
"all_availability_zones" = tobool(null)
"exclude_names" = toset(null) /* of string */
"exclude_zone_ids" = toset(null) /* of string */
"filter" = toset(null) /* of object */
"group_names" = toset([
"ca-central-1",
])
"id" = "ca-central-1"
"names" = tolist([
"ca-central-1a",
"ca-central-1b",
"ca-central-1d",
])
"state" = "available"
"zone_ids" = tolist([
"cac1-az1",
"cac1-az2",
"cac1-az4",
])
}
az0 = "ca-central-1a"
sg_info = "sg-01f056462eff075ab"
subnet_info = "subnet-069d3e3c351d04c49"
vpc_info = "vpc-08135b0d62921449e"

```
MINGW64:/c/Users/Vaishal/Documents/LambtonProject
  - sg_info     = "sg-01f056462eff075ab" -> null
  - subnet_info = "subnet-069d3e3c351d04c49" -> null
  - vpc_info    = "vpc-08135b0d62921449e" -> null

Do you really want to destroy all resources?
  Terraform will destroy all your managed infrastructure, as shown above.
  There is no undo. Only 'yes' will be accepted to confirm.

  Enter a value: yes

aws_route_table_association.main-public[1]: Destroying... [id=rtbassoc-03e98a65916035699]
aws_eip.ec2: Destroying... [id=eipalloc-0156fe767c923beec]
aws_route_table_association.main-public[2]: Destroying... [id=rtbassoc-022177d82dbeba1d9]
aws_route_table_association.main-public[0]: Destroying... [id=rtbassoc-07c84a82514e8511d]
aws_subnet.main-private[0]: Destroying... [id=subnet-0750433a43bbc779e]
aws_subnet.main-private[2]: Destroying... [id=subnet-0d694937bfcf92754]
aws_subnet.main-private[1]: Destroying... [id=subnet-0a98e66c057fc7a2c]
aws_route_table_association.main-public[1]: Destruction complete after 1s
aws_route_table_association.main-public[2]: Destruction complete after 1s
aws_route_table_association.main-public[0]: Destruction complete after 1s
aws_route_table.main-public: Destroying... [id=rtb-0b1532e8e86ce3ec8]
aws_subnet.main-private[2]: Destruction complete after 1s
aws_subnet.main-private[1]: Destruction complete after 1s
aws_subnet.main-private[0]: Destruction complete after 1s
aws_route_table.main-public: Destruction complete after 1s
aws_internet_gateway.main-gw: Destroying... [id=igw-05f34a2f087b6a710]
aws_eip.ec2: Destruction complete after 2s
aws_instance.main: Destroying... [id=i-026421a9931605865]
aws_internet_gateway.main-gw: Still destroying... [id=igw-05f34a2f087b6a710, 10s elapsed]
aws_instance.main: Still destroying... [id=i-026421a9931605865, 10s elapsed]
aws_internet_gateway.main-gw: Still destroying... [id=igw-05f34a2f087b6a710, 20s elapsed]
aws_instance.main: Still destroying... [id=i-026421a9931605865, 20s elapsed]
aws_internet_gateway.main-gw: Destruction complete after 30s
aws_instance.main: Still destroying... [id=i-026421a9931605865, 30s elapsed]
aws_instance.main: Destruction complete after 32s
aws_subnet.main-public[1]: Destroying... [id=subnet-0733d55d7956d432e]
aws_subnet.main-public[2]: Destroying... [id=subnet-06092e7be48e888ce]
aws_subnet.main-public[0]: Destroying... [id=subnet-069d3e3c351d04c49]
aws_security_group.myinstance: Destroying... [id=sg-01f056462eff075ab]
aws_security_group.myinstance: Destruction complete after 1s
aws_subnet.main-public[2]: Destruction complete after 1s
aws_subnet.main-public[1]: Destruction complete after 1s
aws_subnet.main-public[0]: Destruction complete after 1s
aws_vpc.main: Destroying... [id=vpc-08135b0d62921449e]
aws_vpc.main: Destruction complete after 1s

Destroy complete! Resources: 15 destroyed.
```

AWS CLOUD:

All resources are deployed on the AWS cloud via Terraform. Resources used in this demo are: VPC, Subnets public/private, Route-table, Internet Gateway, Security Group, NACL, Elastic IP, EBS, Key Pair, EC2 instance. EC2 instance of type T3.medium is used and the flavor is Ubuntu 20.04.

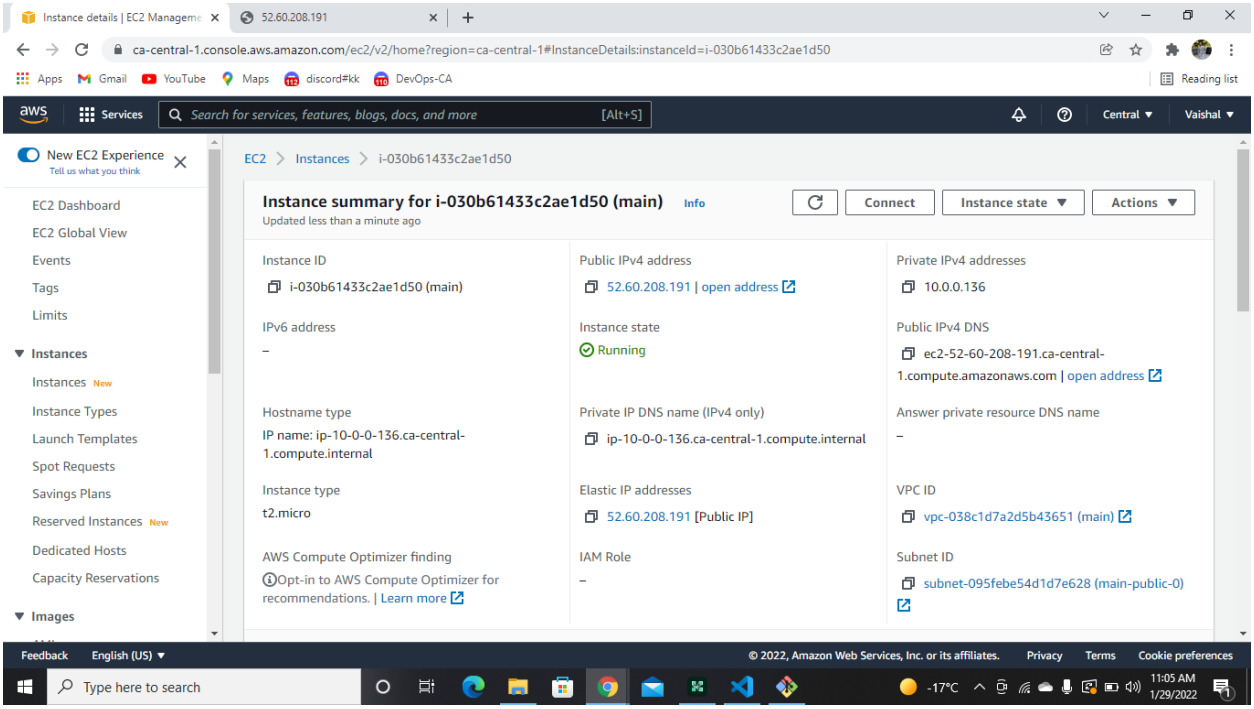

DOCKER:

Docker, being a containerization technology, I have used it to containerize the whole ELK stack along with Nginx. Nginx is a web server whose logs are collected and stored on Elasticsearch. These logs are procured by Kibana which in turn helps us to visualize the data.

The stack is deployed using docker and docker-compose. Docker is used to make the application/ELK stack portable and platform independent.

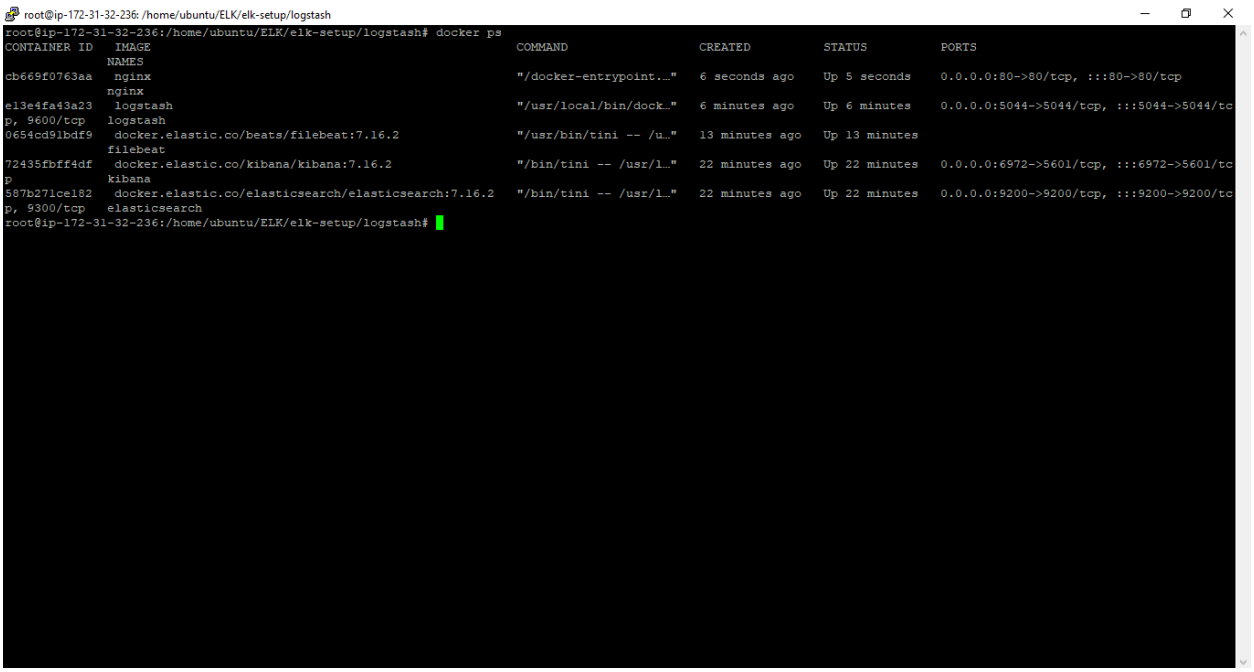

ANSIBLE:

Ansible is used to automate configurations of the ELK stack. I have written an Ansible playbook using YAML. So, we can run playbooks to multiple servers if needed, without any intervention into the virtual machine remotely. For now, I have written a playbook for installing docker on a remote machine. In near future, I will write an entire ELK configuration playbook that can be deployed remotely using the Ansible playbook.

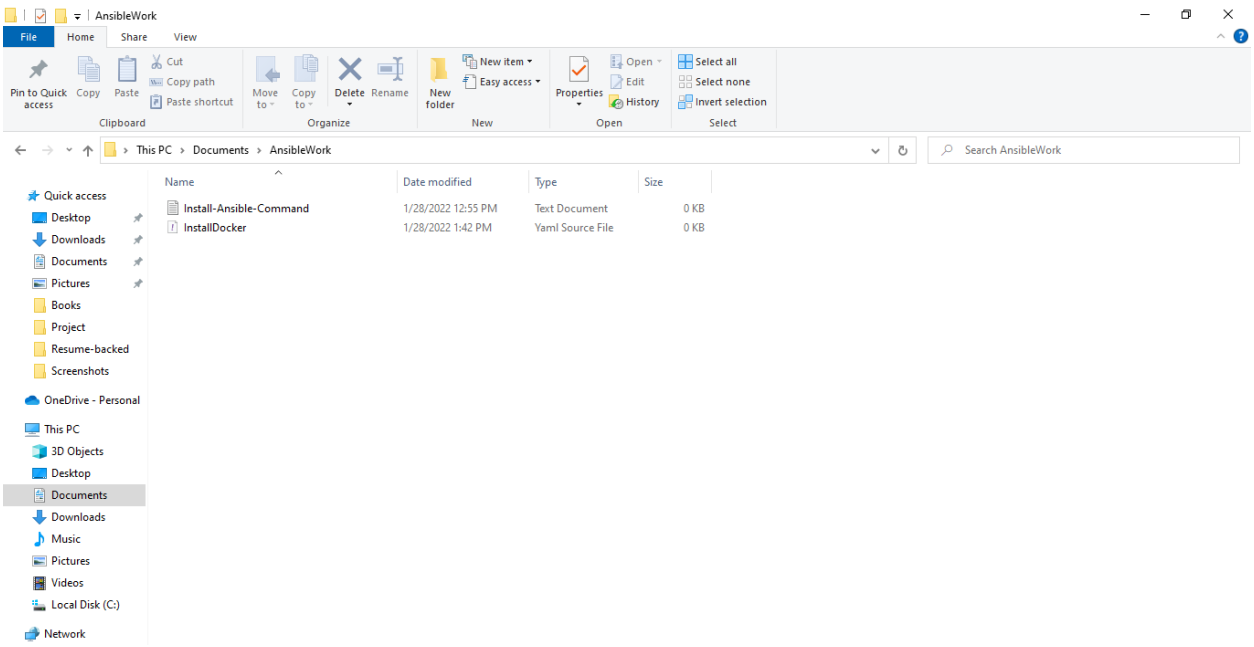

ELK STACK:

It stands for Elastic, Logstash, Kibana. The stack is deployed to make a one-stop self-hosted solution for monitoring, and logging management systems.

In this demo, Filebeat collects the logs of various docker containers and sends them out to Logstash, where filtering/indexing of the logs takes place and then Logstash sends it to Elasticsearch. All the data gets stored in Elasticsearch. From Elasticsearch, Kibana reads the data and makes it available for visualizing the logs as per our needs.

The attached screenshots show the Elasticsearch workspace, and the index management enabled on Kibana to process and read the logs from Elasticsearch. Lastly, the Kibana dashboard shows where data/logs can be visualized.

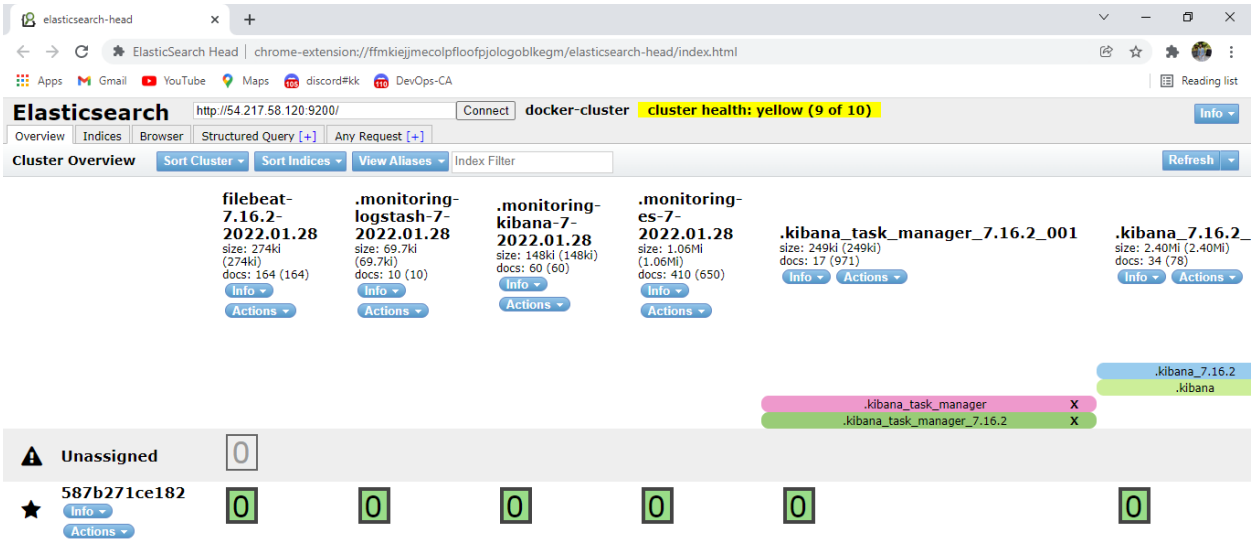
elasticsearch-head
ElasticSearch Head | chrome-extension://ffmkiejjmecolpfloofpjologoblkegm/elasticsearch-head/index.html
Apps Gmail YouTube Maps discord#kk DevOps-CA Reading list
Elasticsearch
http://54.217.58.120:9200/
Connect
docker-cluster
cluster health: yellow (9 of 10)
Info
Overview Indices Browser Structured Query [+] Any Request [+]
Cluster Overview
Sort Cluster Sort Indices View Aliases Index Filter
Refresh
filebeat-7.16.2-2022.01.28
size: 274ki (274ki)
docs: 164 (164)
.monitoring-logstash-7-2022.01.28
size: 69.7ki (69.7ki)
docs: 10 (10)
.monitoring-kibana-7-2022.01.28
size: 148ki (148ki)
docs: 60 (60)
.monitoring-es-7-2022.01.28
size: 1.06Mi (1.06Mi)
docs: 410 (650)
.kibana_task_manager_7.16.2_001
size: 249ki (249ki)
docs: 17 (971)
.kibana_7.16.2_
size: 2.40Mi (2.40Mi)
docs: 34 (78)
Info
Actions
.kibana_7.16.2
.kibana
.kibana_task_manager
.kibana_task_manager_7.16.2
Unassigned
0
587b271ce182
Info
Actions
0

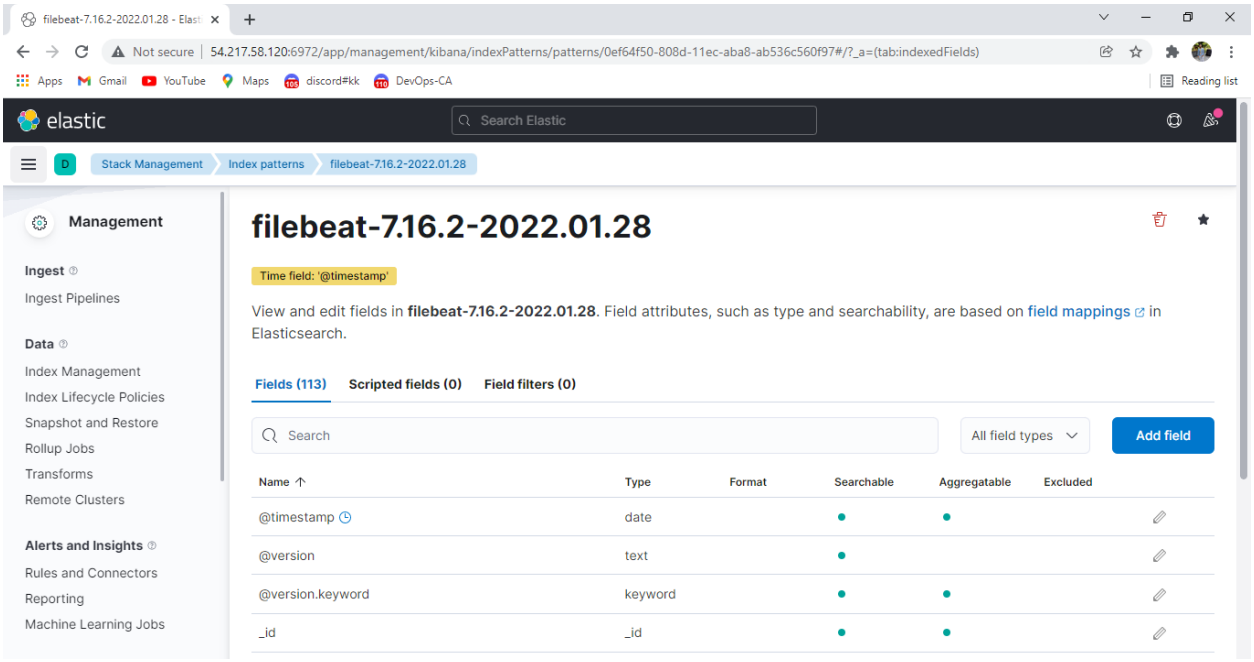
filebeat-7.16.2-2022.01.28 - Elast
Not secure | 54.217.58.120:6972/app/management/kibana/indexPatterns/patterns/0ef64f50-808d-11ec-aba8-ab536c560f97#/?_a=(tab:indexedFields)
Apps Gmail YouTube Maps discord#kk DevOps-CA Reading list
elastic
Search Elastic
Stack Management
Index patterns
filebeat-7.16.2-2022.01.28
Management
Ingest
Ingest Pipelines
Data
Index Management
Index Lifecycle Policies
Snapshot and Restore
Rollup Jobs
Transforms
Remote Clusters
Alerts and Insights
Rules and Connectors
Reporting
Machine Learning Jobs
filebeat-7.16.2-2022.01.28
Time field: '@timestamp'
View and edit fields in filebeat-7.16.2-2022.01.28. Field attributes, such as type and searchability, are based on field mappings in Elasticsearch.
Fields (113)
Scripted fields (0)
Field filters (0)
Search
All field types
Add field
Name
Type
Format
Searchable
Aggregatable
Excluded
@timestamp
date
@version
text
@version.keyword
keyword
_id
_id

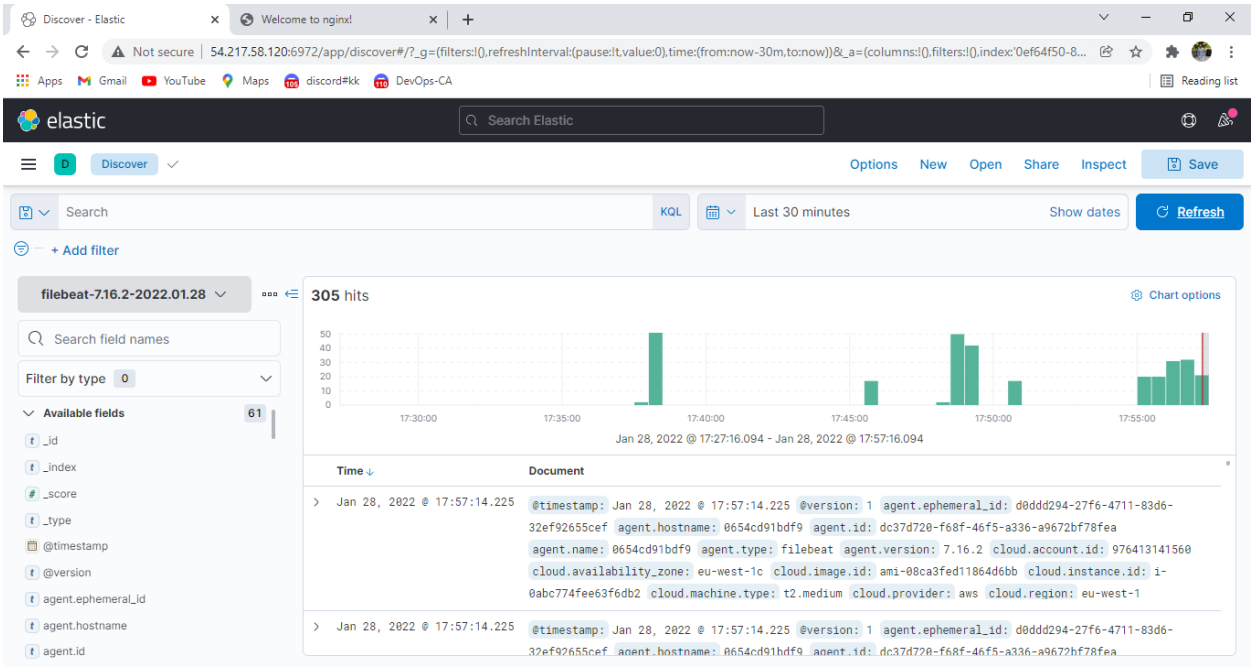

CI/CD:

I will make it CI/CD enabled using either GitHub actions or Jenkins. It is the future plan, which will be ready soon. Which will hold the terraform code. So, infrastructure will be automatically deployed on AWS with just a simple git push command.

REFERENCE:

Blog written by me: https://faun.pub/elk-elasticsearch-logstash-kibana-conceptual-tutorial-for-beginners-2a7a827305b8